

Testing of several distributed file-system (HadoopFS, CEPH and GlusterFS) for supporting the HEP experiments analisys.

Giacinto DONVITO
INFN-Bari

Agenda

- Introduction on the objective of the test activities
- HadoopFS
- GlusterFS
- CEPH
- Tests and results
- Conclusion and future works

Introduction on the objective of the test activities

- The aim of the activity is to verify:
 - Performance
 - Reliability
 - Features
- Considering solutions that provides software redundancy
 - A site could use commodity hardware to achieving high level of data availability
- The scalability should be guaranteed at the order of few PetaByte

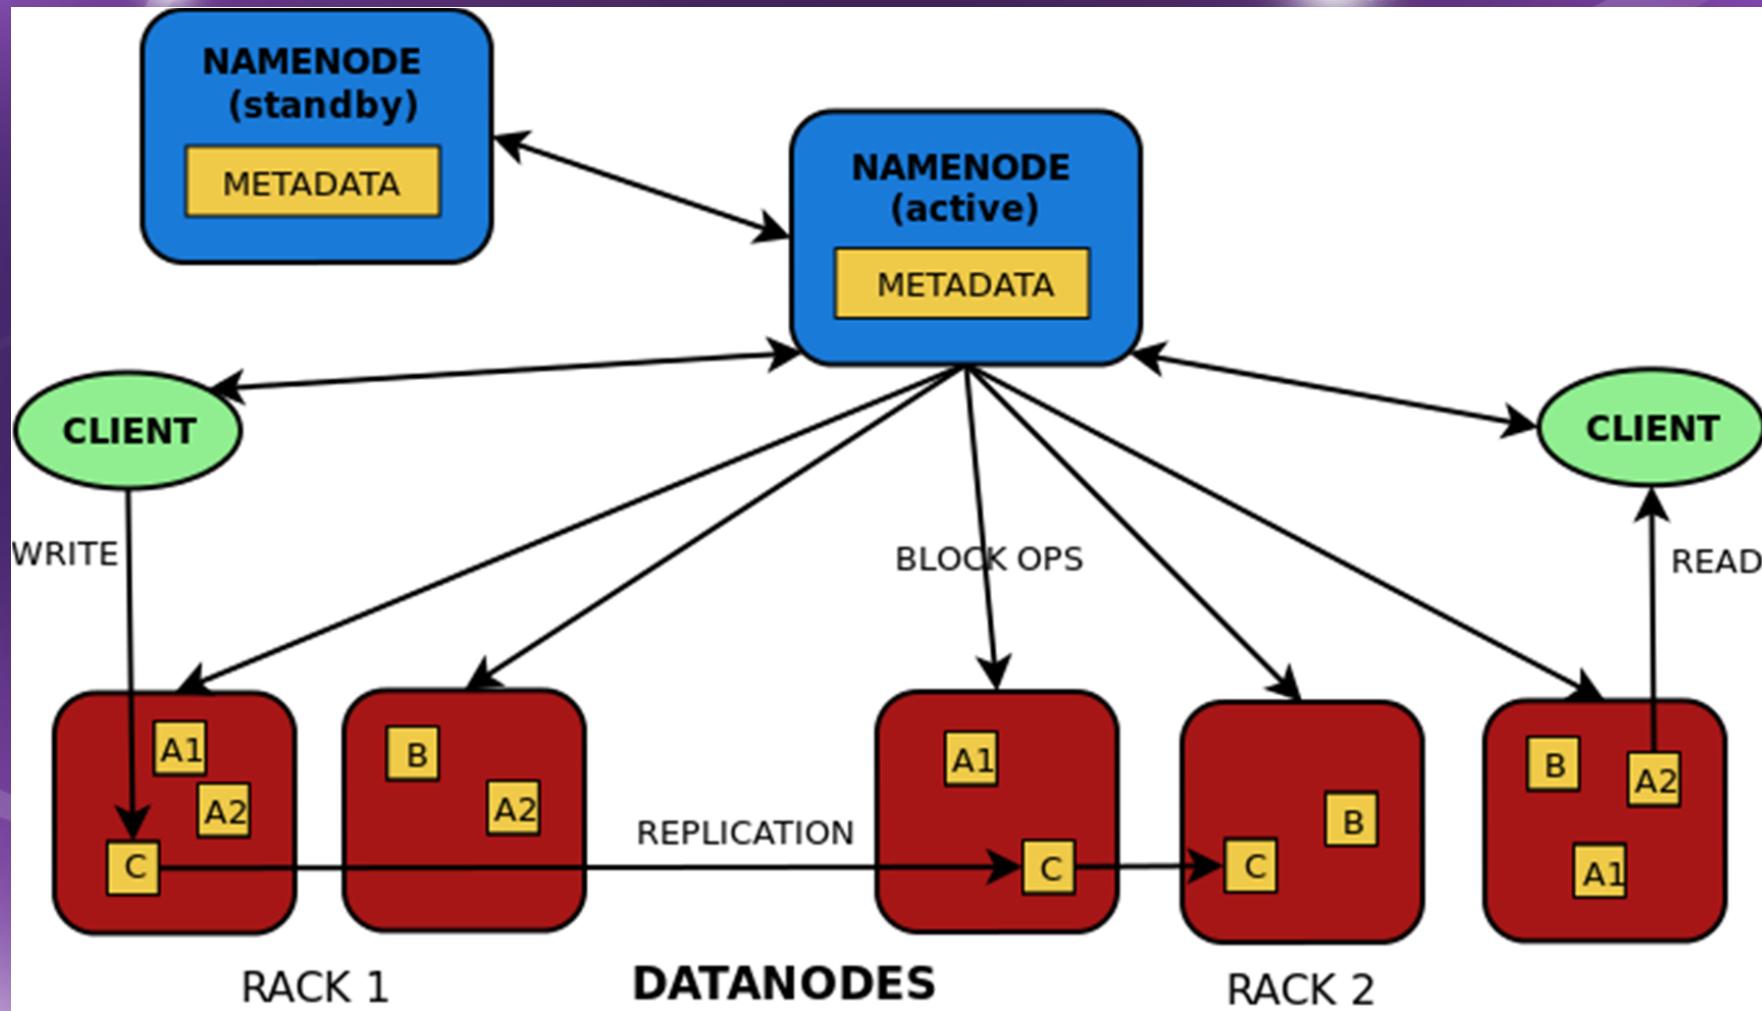
Introduction on the objective of the test activities

- The focus is to serve typical Tier2/Tier3 sites for LHC experiments
 - Supporting interactive usage
 - Running data analysis
 - Supporting SRM, gridftp, Xrootd
 - Being prepared for the new cloud storage techniques
- Open Source solutions

HadoopFS

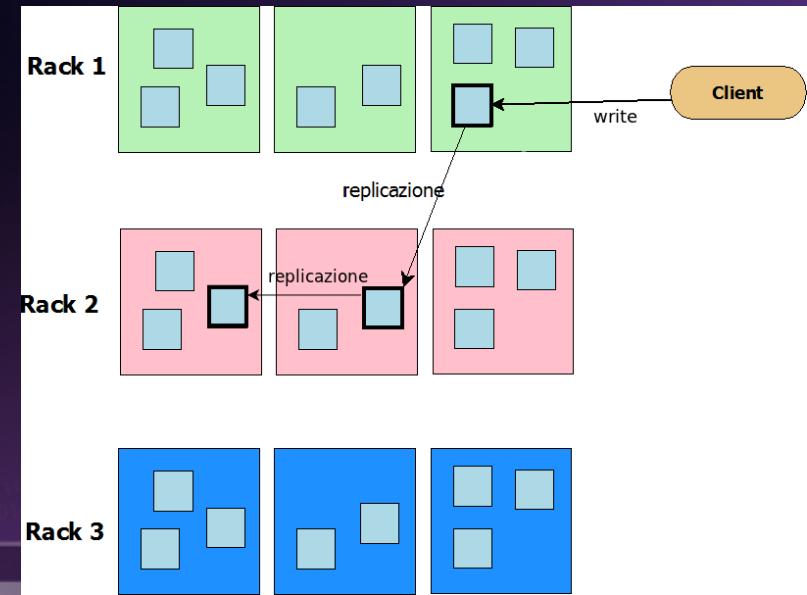
- Apache Hadoop Distributed File System:
 - Open-source
 - Developed in Java
 - Large dataset
 - Fault tolerant
 - Commodity hardware
 - High throughput
 - Scalable
 - Rack awareness

HadoopFS



HadoopFS

- "The primary objective of HDFS is to store data reliably even in the presence of failures" (Hadoop documentation)
 - File are split in chunk (default 64MB)
 - `dfs.blocksize`
 - Placement policy (default):
 - 3 replicas
 - 1 replica in the local rack
 - 2 replicas in the remote rack
 - `dfs.replication`

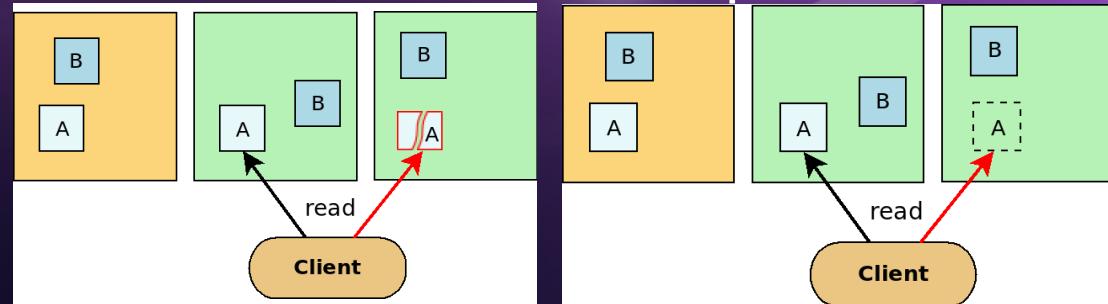


HadoopFS functionality test

- We have executed several test to check the behavior at several types of failures:
 - Metadata failures
 - Client retry, active-standby namenode
 - Datanode failures:
 - During write operation, during read operation, in case of data corruption, mis-replicated blocks, under and over replicated blocks
- We always succeeded to fulfill the expected behavior and no (un-expected) data-loss were registered

HadoopFS

Data Corruption



NameNode 'pccms61.ba.infn.it:9000' (active)

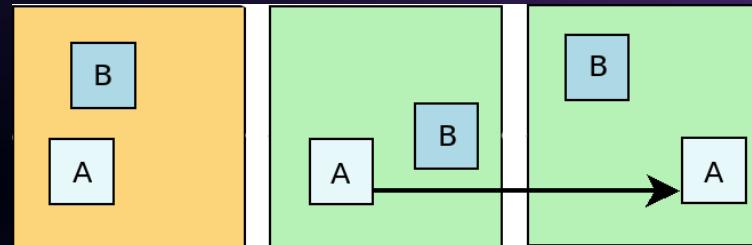
Started: Thu Aug 22 11:45:05 CEST 2013
Version: 2.0.0-cdh4.1.1, 581959ba23e4af85af8db98b7687662fe9c5f20
Compiled: Tue Oct 16 10:39:59 PDT 2012 by jenkins from Unknown
Upgrades: There are no upgrades in progress.
Cluster ID: CID-9b734b6d-3611-4eb7-ab5e-49419f75dc3a
Block Pool ID: BP-1130807058-212.189.205.51-1340275038748

[Browse the filesystem](#)
[NameNode Logs](#)

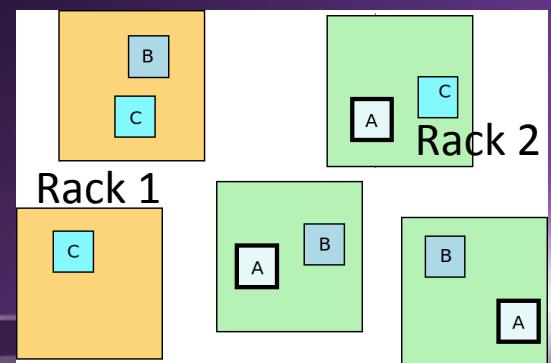
Cluster Summary

Security is OFF
 719577 files and directories, 798359 blocks = 1517936 total.
 Heap Memory used 3.29 GB is 69% of Committed Heap Memory 4.71 GB. Max Heap Memory is 8.89 GB.
 Non Heap Memory used 42.98 MB is 71% of Committed Non Heap Memory 60.38 MB. Max Non Heap Memory is 130 MB.
WARNING : There are 1642 missing blocks. Please check the logs or run fsck in order to identify the missing blocks.
 See the Hadoop FAQ for common causes and potential solutions.

| | | |
|-----------------------------------|---|------------------------------|
| Configured Capacity | : | 135.31 TB |
| DFS Used | : | 23.48 TB |
| Non DFS Used | : | 6.66 TB |
| DFS Remaining | : | 105.16 TB |
| DFS Used% | : | 17.36 % |
| DFS Remaining% | : | 77.72 % |
| Block Pool Used | : | 23.48 TB |
| Block Pool Used% | : | 17.36 % |
| DataNodes usages | : | Min % Median % Max % stdev % |
| | | 0 % 25.59 % 100 % 32.31 % |
| Live Nodes | : | 155 (Decommissioned: 9) |
| Dead Nodes | : | 97 (Decommissioned: 36) |
| Decommissioning Nodes | : | 0 |
| Number of Under-Replicated Blocks | : | 77 |

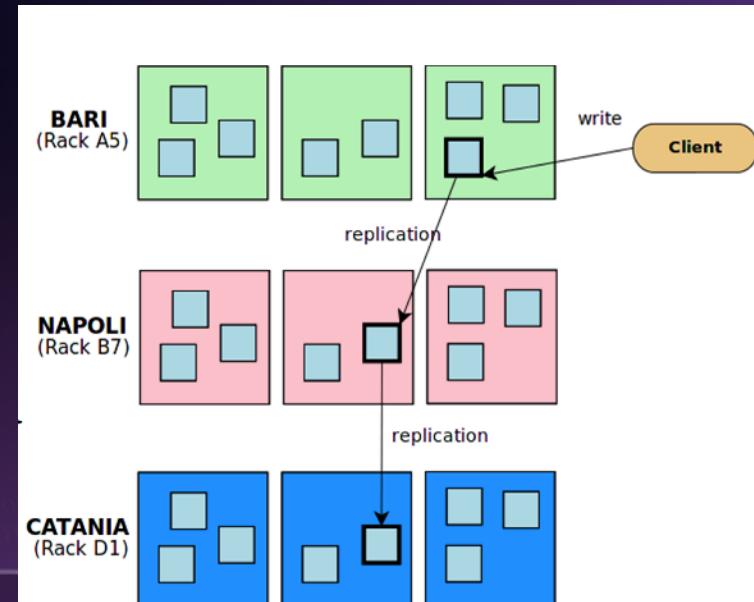


Mis-replicated blocks



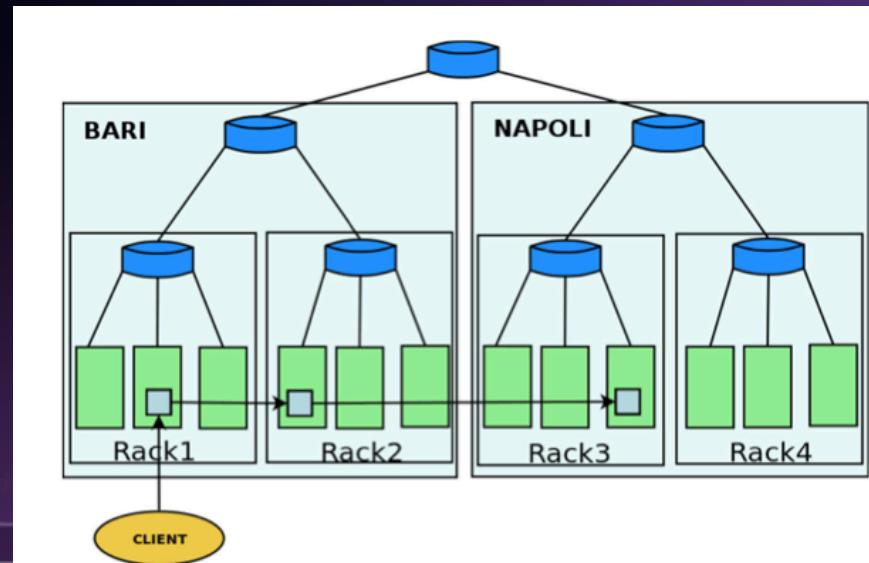
HadoopFS: our development

- One Replica Policy
 - 1 replica per rack
 - Increasing the reliability
 - Increasing the available bandwidth for read operation



HadoopFS: our development

- Hierarchical Policy
 - It is able to exploit a geographically distributed infrastructure
 - 2 replicas in the source site in 2 different racks
- The data will survive also to the loss of a complete site



HadoopFS pros&cons

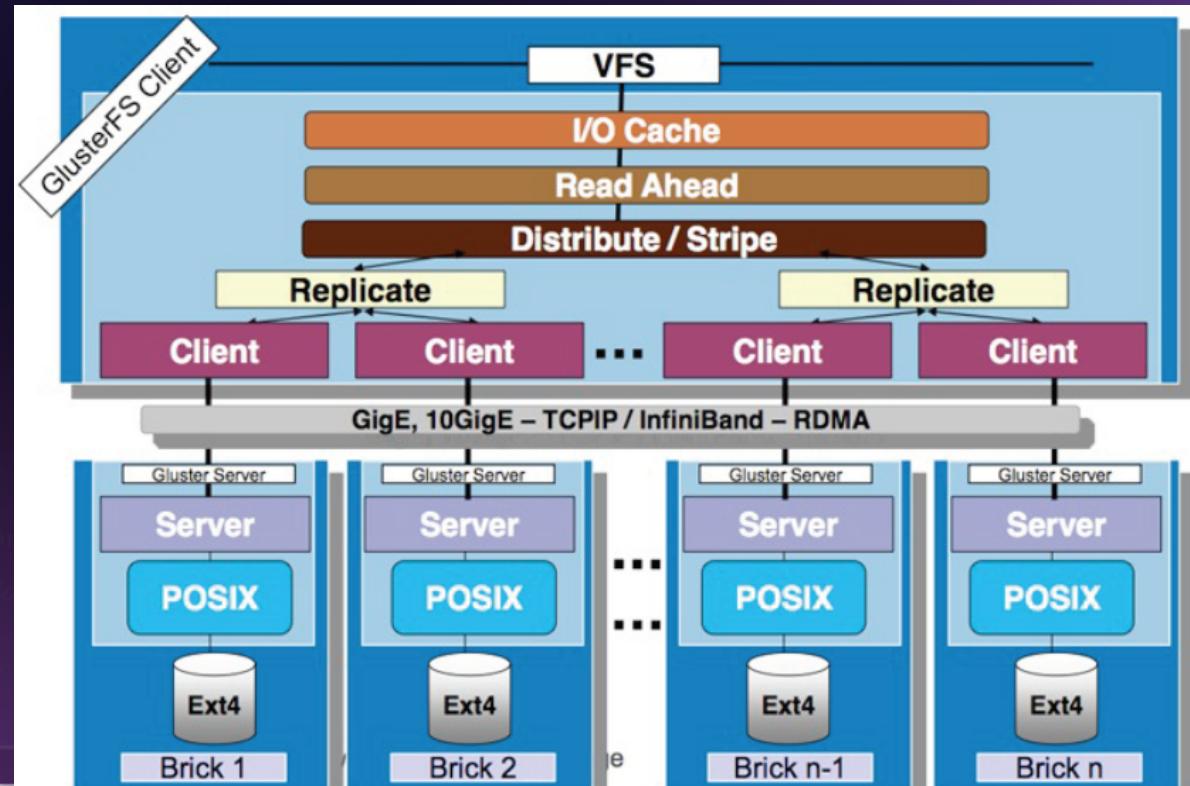
- MapReduce
- Dynamic self-healing
- Great Scalability
 - Already tested in few big Tier2 in LHC and many companies
- Web monitoring interface
- Support for SRM (Bestman) and gridftp/xrootd (Nebraska)
- Non strictly-posix compliance
 - Fuse based
- No support for new cloud storage technologies

GlusterFS

- OpenSource solution acquired by RedHat
- Could be used both with disk in the WN and with standard infrastructures based on disk servers (SAN/DAS)
- Written in C under GPLv3
- Posix compliance
- Exploit NFS protocol
- Available on many platforms (RedHat, Debian, MacOS, NetBSD, OpenSolaris)
- Support also new storage cloud technologies (Block Storage, Object Storage, etc)
 - It is based on Swift (OpenSource Object Storage developed within OpenStack framework)

GlusterFS

- Working behavior:
 - The client exploit a FUSE module to access file and implement advanced policy (Distribute/Stripe/Replica, etc)
- The client and server could exploit both TCP/IP and infiniband connections
- The server hosts data on standard file-systems (ext4, xfs, etc)



GlusterFS

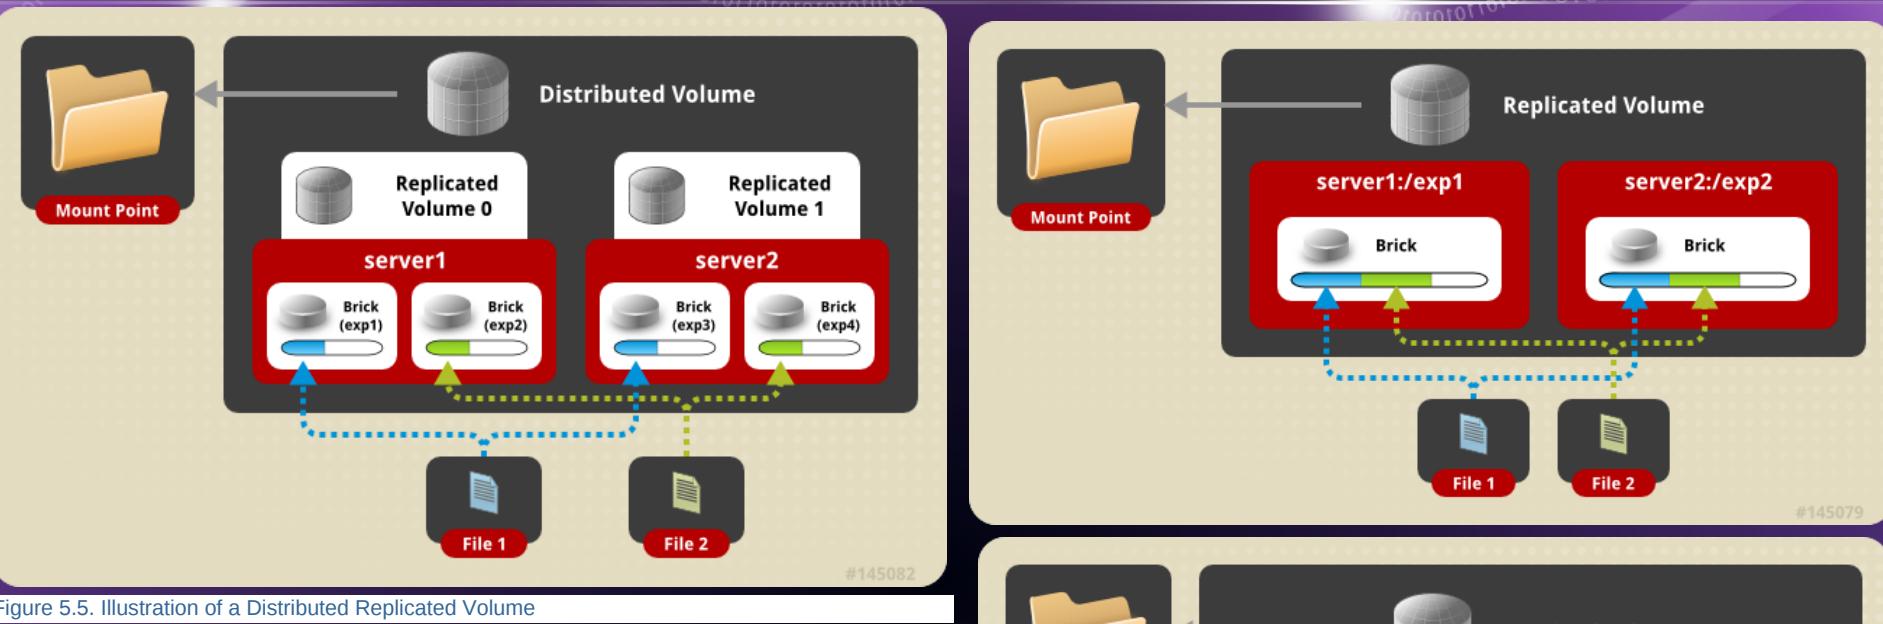


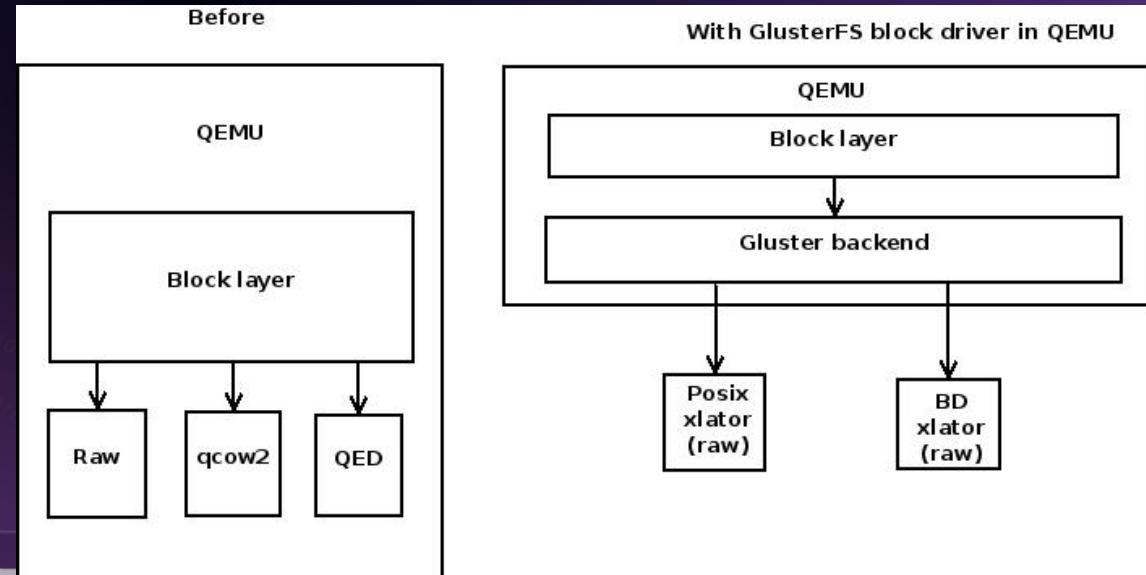
Figure 5.5. Illustration of a Distributed Replicated Volume

- Working behavior:
 - Striped Volume
 - Replicated Volume
 - Distributed Volume
 - Striped Replicated Volume
 - Distributed Replicated Volume

Figure 5.6. Illustration of a Striped Replicated Volume

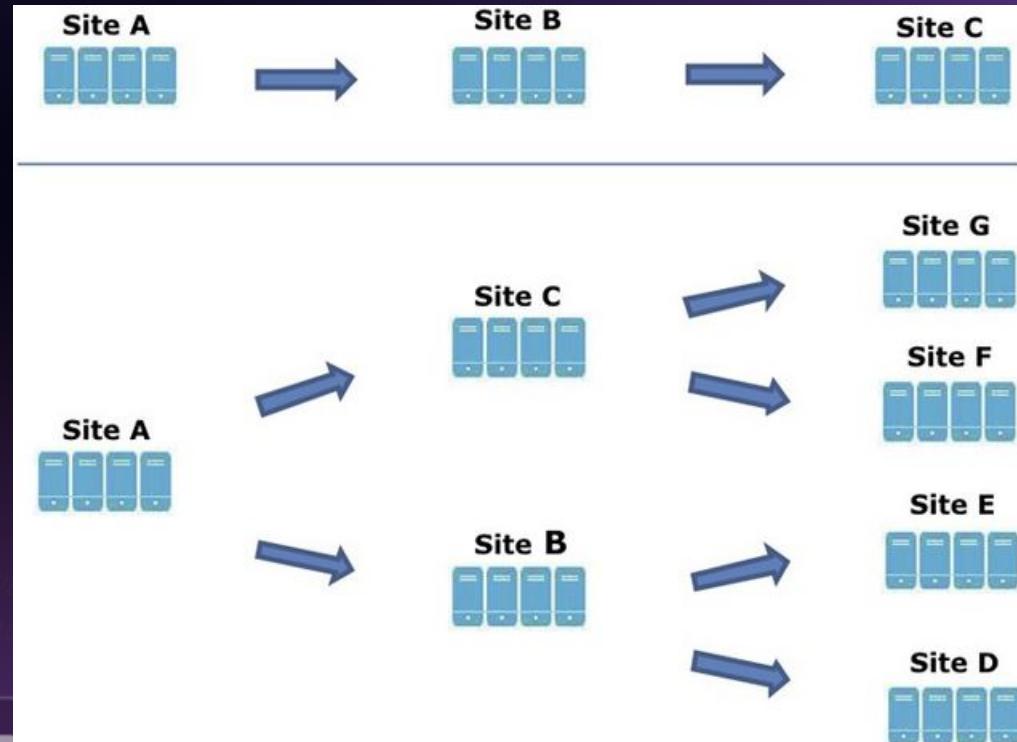
GlusterFS

- POSIX ACL support over NFSv3
- Virtual Machine Image Storage
 - qemu – libgfapi integration
 - improvements in performance for VM image hosting
- Synchronous replication improvements
- Distributed-replicated and Striped-replicated are very important in the contest where performance and data availability is important



GlusterFS

- GlusterFS provides a geographical replication solution
 - Could be useful as disaster recovery solution
 - It is based on the paradigm of active-backup
 - It is based on rsync
 - It is possible to replicate the whole file-system or a part of it
 - It could be used also from one site to multiple instances of GlusterFS on different sites



Glusterfs pros&cons

- Easy to install and configure
- Fully posix compliance
- Many available configuration
- Great performance
- Provides interesting cloud storage solutions
- Some instabilities and data loss in some specific situations
- There are no many scalability reports beyond petabyte

CEPH file-system

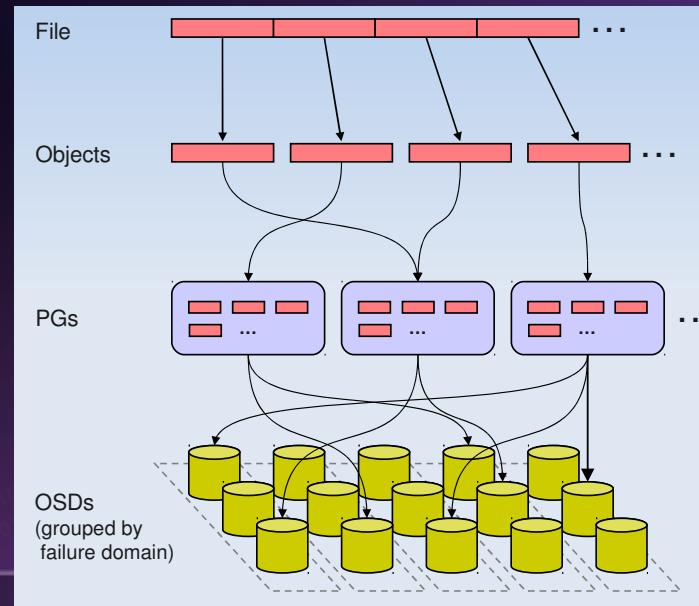
- Development started in 2009
- Now it is acquired by a company (Inktank) also if it is still an completely OpenSource projects
- It is integrated by default in the Linux Kernel since 2.6.34 release (may 2010)
- It could use, although not already at “production level”, BTRFS (B-tree file system) as backend
 - Several interesting features (Raid0/1, and soon Raid5/6, data deduplication, etc) implemented at software level

CEPH file-system

- Designed to be scalable and fault-tolerant
 - In order to support >10'000 disk server
 - Up to 128 metadata server (could serve up to 250kops/s aggregate)
- CEPH can provide three different storage interfaces: Posix (both at kernel level and using fuse), Block and Object storage
- Several IaaS cloud platforms (i.e.: OpenStack, CloudStack) officially supports CEPH to provides Block Storage solution
- The suggested configuration do not require/suggest the use of any hardware raid: the data availability is implemented at software level

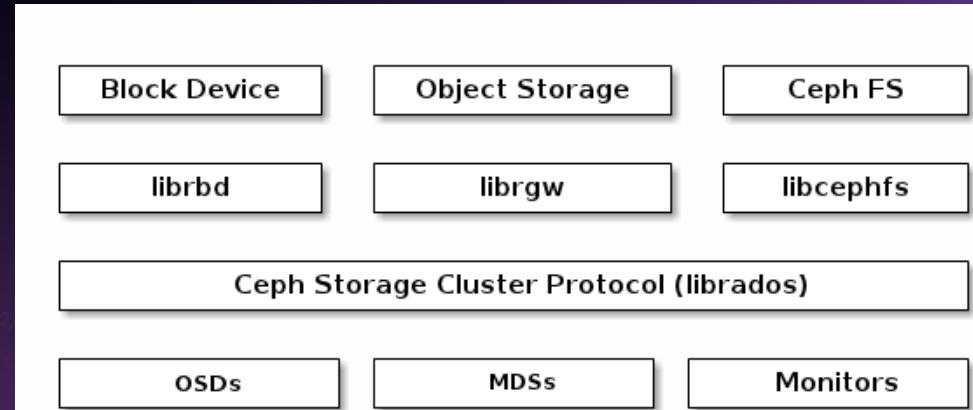
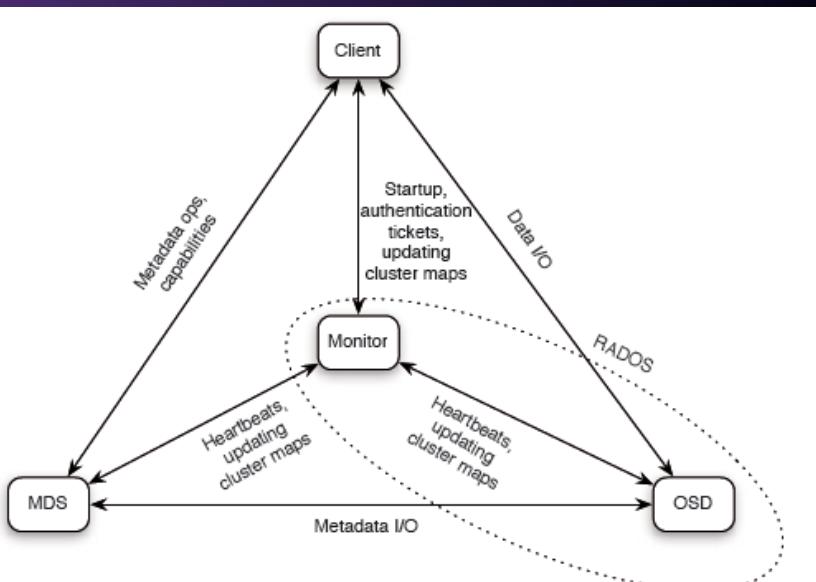
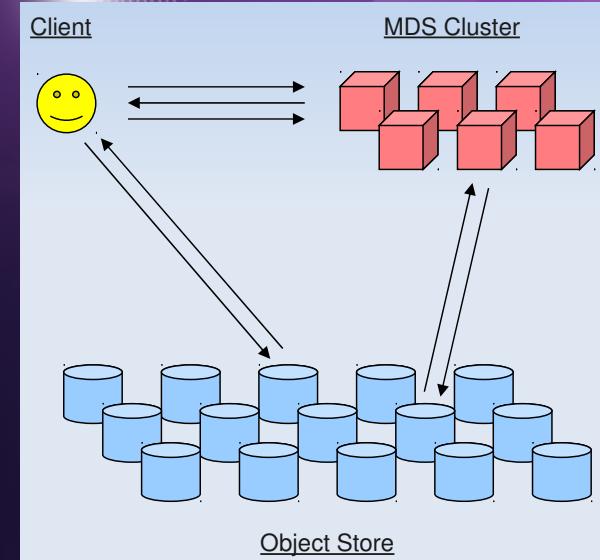
CEPH file-system

- The data distribution is based on an hash function
 - No query needed to know the location of a given file
- This means that the mapping is “unstable”:
 - Adding a disk server, mean that the whole cluster need to reshuffling the location of the data
- It is possible to define “failure domain” at the level of: disk, server, rack
- Data placement rules could be customized:
 - “tre different copies of the same file in three different racks”
- All the datanodes knows the exact location of all the files in the cluster



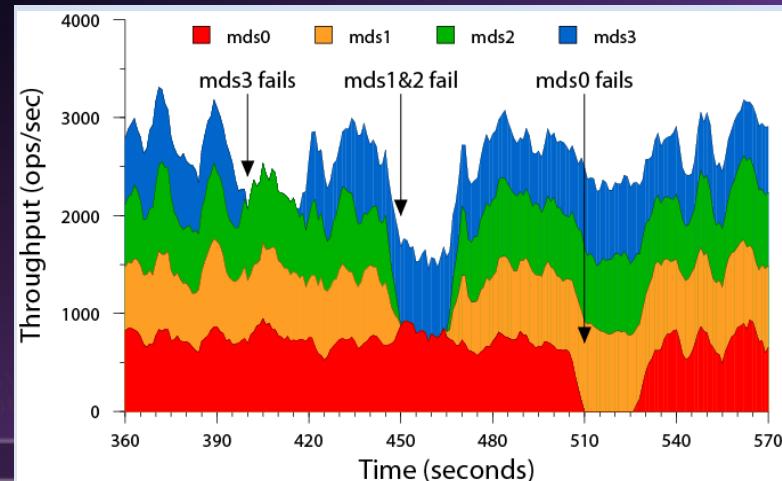
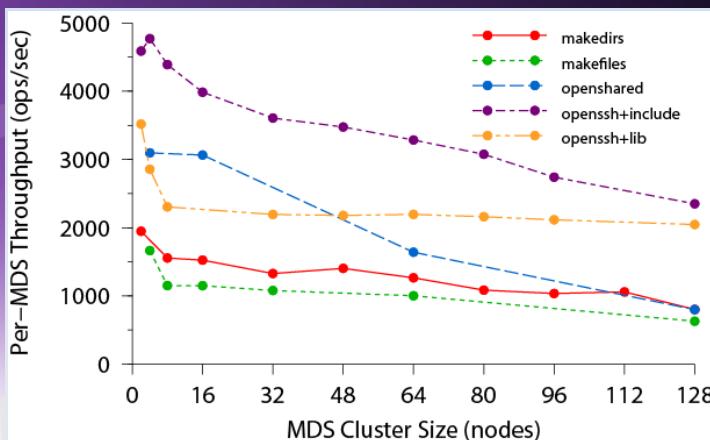
CEPH file-system

- Monitor: manages the heartbeats among nodes
- MDS: manages I/O on metadata
- OSD: contains the objects
- The client will interact with all the three services
- A 10 node cluster will be composed by:
 - 3 monitor node
 - 3 MDS node
 - 7 OSD node



CEPH file-system

- The three storage interfaces (posix, block and object) are different gateways on the same objects APIs
- The object could be stored also “striped” in order to increase the performances
 - Object Size, Stripe Width, Stripe Count
- Data Scrubbing: it is possible to periodically check the data consistency (to avoid inconsistencies between data and metadata, and or data corruptions)



CEPH functionalities test

- The “quorum” concept is used for each critical service (there should be odds numbers of instances):
 - If 2 over 3 services are active the client could read and write. Is only one is active the client could only read
- We verified the behaviour in case of failure of each service:
 - The High Availability worked always as expected
 - We tested both failure in data and metadata services
 - Both using posix and RBD interfaces
- We tested also the possibility to export the storage using standard NFS protocols
 - It works quite well both using RBD and POSIX interface
 - Was very unstable using kernel interface

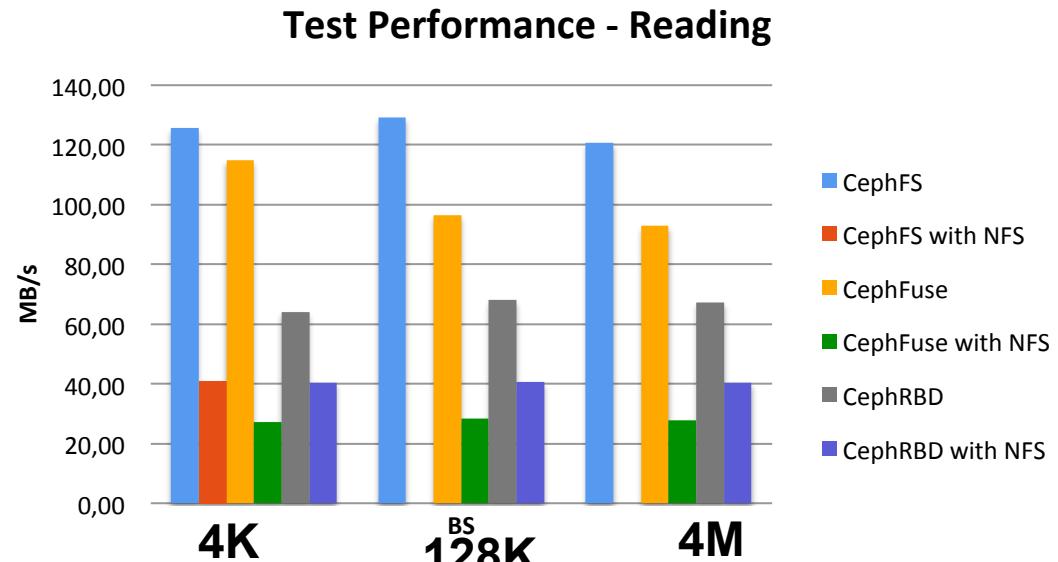
CEPH RBD

- CEPH RBD features:
 - Thinly provisioned
 - Resizable images
 - Image import/export
 - Image copy or rename
 - Read-only snapshots
 - Revert to snapshots
 - Ability to mount with Linux or QEMU KVM clients
- In OpenStack it is possible to use CEPH both as device in Cinder (Block storage server) and for hosting virtual images in Glance (Image Service)
- CEPH provides an Object Storage solution that has interfaces compatible with both S3 (Amazon) and Swift (OpenStack)

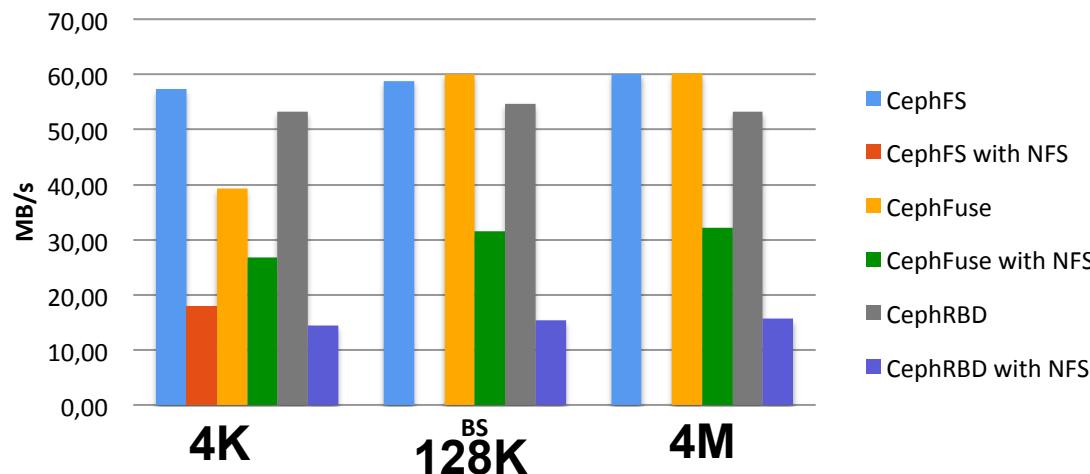
CEPH Performance test

BS

4K
128K
4M



Test Performance - Writing



Virtual Machine

CEPH pros&cons

- Complete storage solution (supports all the storage interfaces: posix, object, block)
- Great scalability
- Fault-tolerant solution
- Difficult to install and configure
- Performance issues
- Some instabilities while under heavy load

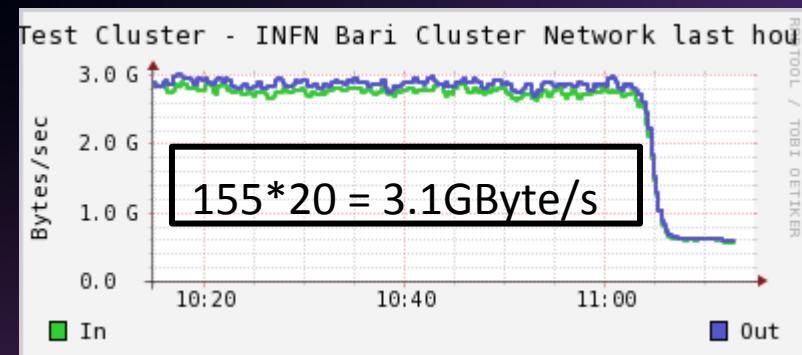
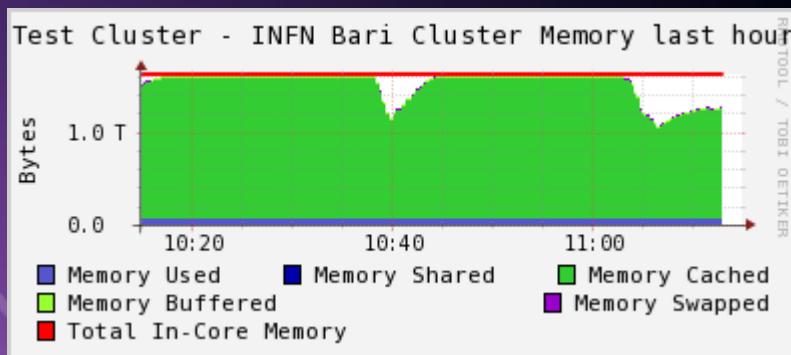
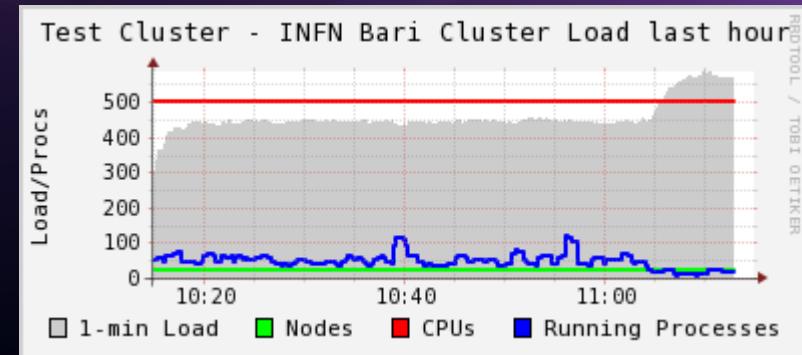
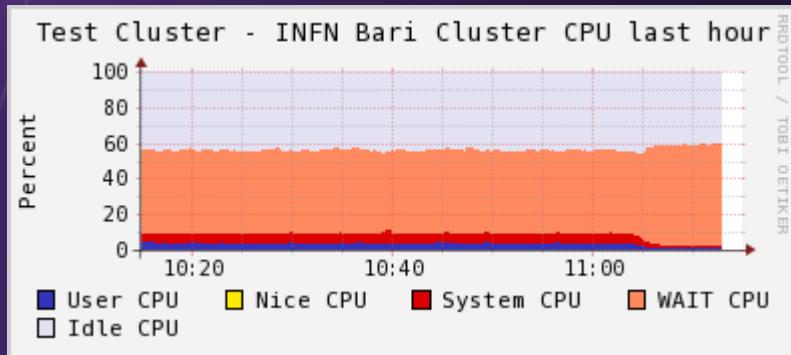
HDFS v2 CDH 4.1.1 (by USCMS) Nebraska)

- 20 datanodes, 1 namenode
- Chunk size: 128MB, Rdbuffer: 128MB, Big_writes active
 - # iozone -r 128k -i 0 -i 1 -i 2 -t 24/36 -s 10G

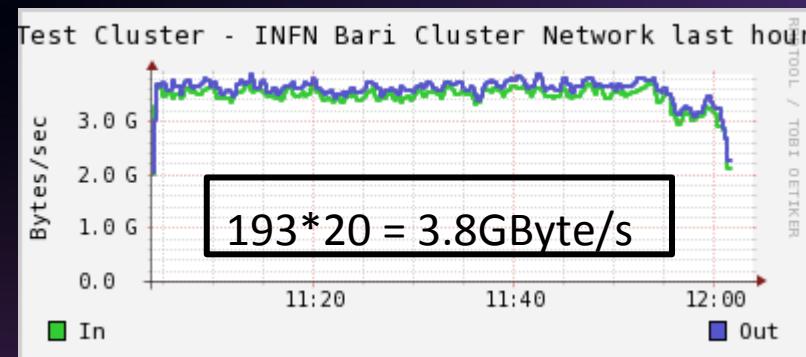
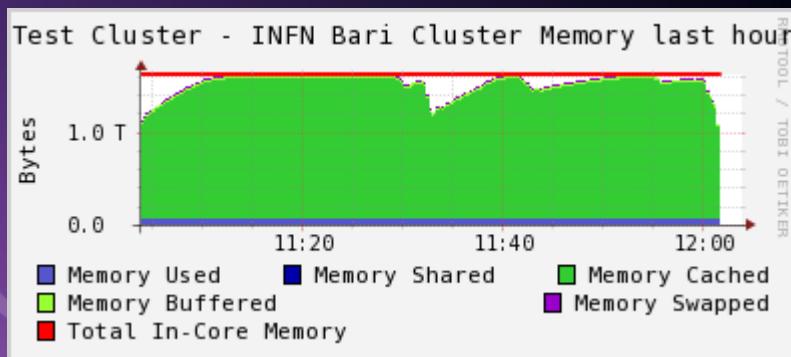
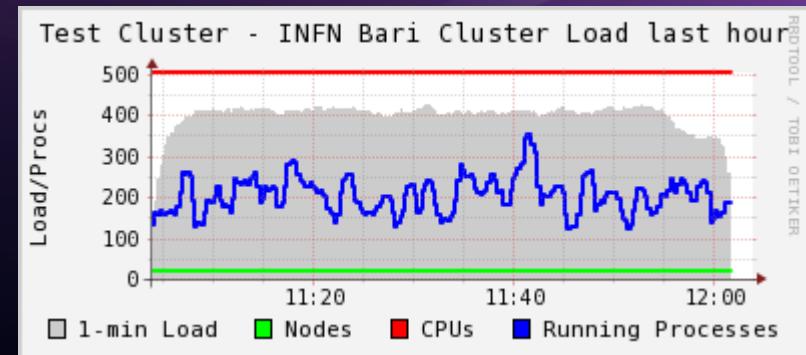
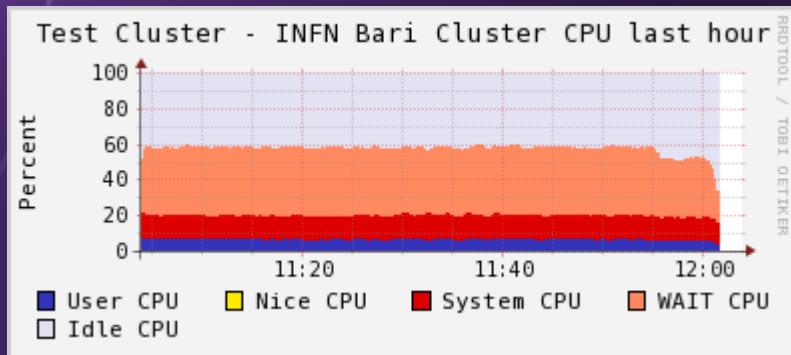
MB/s

| | 24 Threads | 36 Threads |
|---------------|------------|------------|
| Initial Write | 239.72 | |
| Re-write | | |
| Random Write | | |
| Initial Read | 155.18 | 193.65 |
| Re-read | 151.33 | 207.43 |
| Random Read | 29.06 | 39.98 |

HDFS - 24 threads



HDFS - 36 threads



Ceph Cuttlefish (0.61)

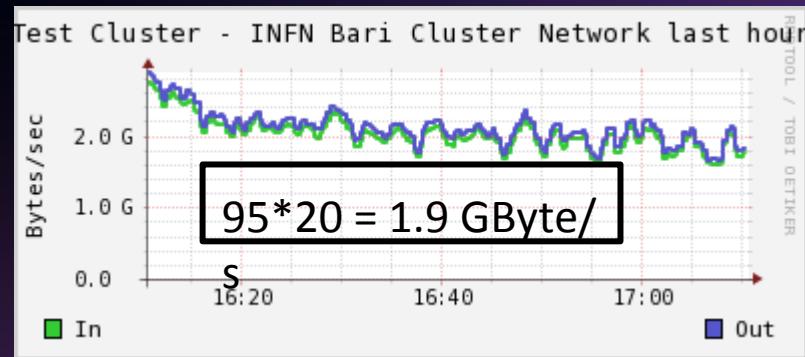
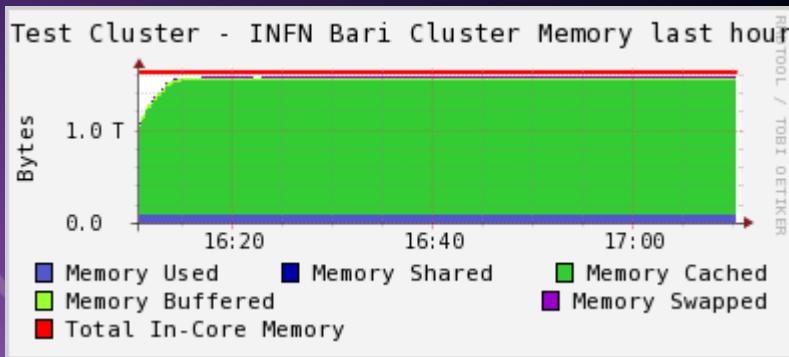
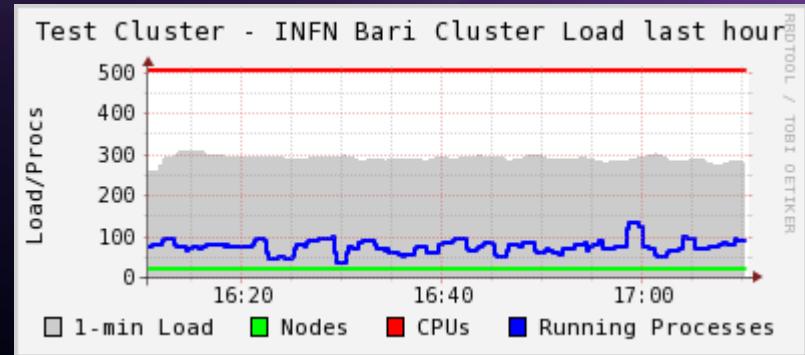
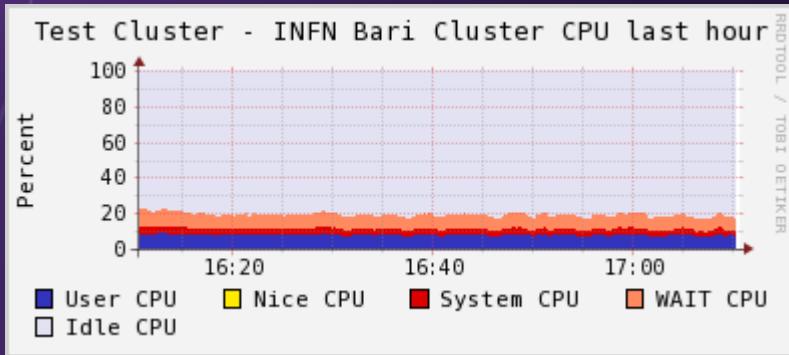
- 3 Mon, 1 Mds, 120 osd (6osd * 20nodi)
- On all the nodes (SLC6)
 - # iozone -r 128k -i 0 -i 1 -i 2 -t 24 -s 10G

MB/s

| | 24 Threads |
|----------------------|-------------------|
| Initial Write | 52.49 |
| Re-write | 54.05 |
| Random Write | ERROR |
| Read | 95.38 |
| Re-read | 102.04 |
| Random Read | ERROR |

$$95*20 = 1.9 \text{ GByte/s}$$

Ceph Cuttlefish (0.61)



Ceph Dumpling (0.67.3)

- 3 Mon, 1 Mds, 95 osd (5osd * 19nodi)
- On all the nodes (SLC6)
 - # iozone -r 128k -i 0 -i 1 -i 2 -t 24 -s 10G

MB/s

| | 24 Threads |
|----------------------|-------------------|
| Initial Write | 18.93 |
| Re-write | 19.31 |
| Random Write | 13.96 |
| Read | 53.40 |
| Re-read | 57.29 |
| Random Read | 5.13 |

$$53 \times 19 = 1.0 \text{ GByte/s}$$

Ceph-Dev (0.70)

- 3 Mon, 1 Mds, 15 osd (5osd * 3nodi)
- On all the nodes (Ubuntu 12.04)

○ # iozone -r 128k -i 0 -i 1 -i 2 -t 24 -s 10G

MB/s

| | 24 Threads |
|----------------------|-------------------|
| Initial Write | 51,06 |
| Re-write | 60,05 |
| Random Write | 7,00 |
| Read | 101,58 |
| Re-read | 133,61 |
| Random Read | 12,05 |

Gluster v3.3

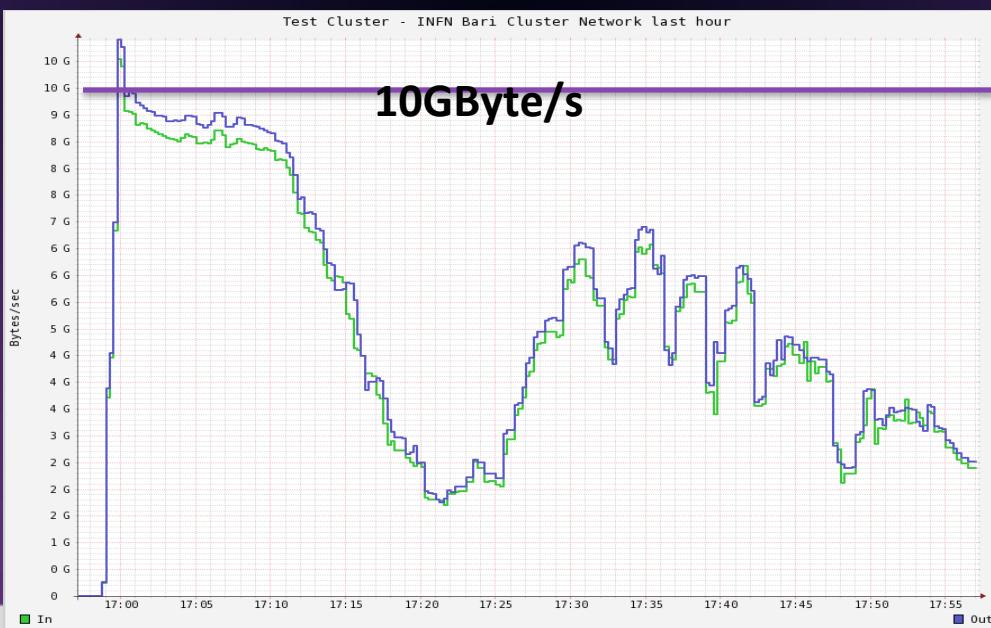
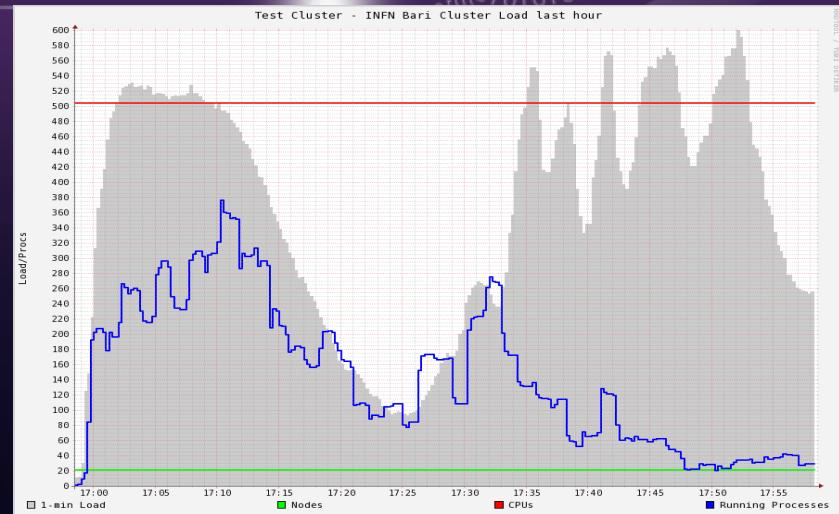
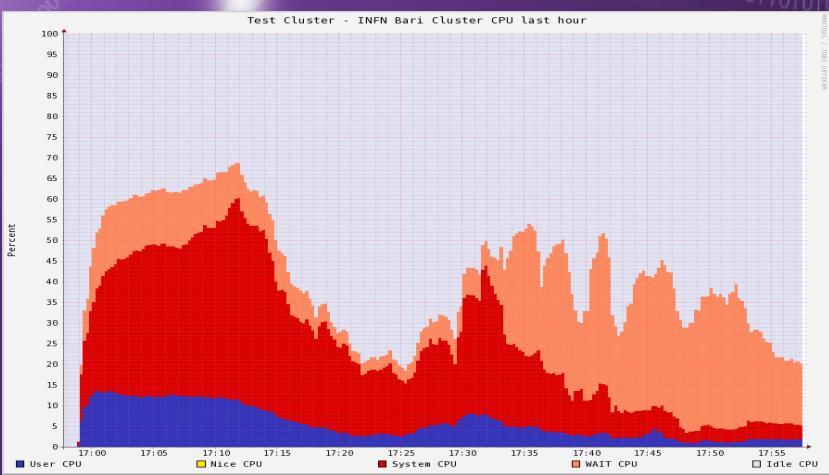
- 21 nodes, 6 brick per node
- On all the nodes (SLC6)
 - # iozone -r 128k -i 0 -i 1 -i 2 -t 24 -s 10G

MB/s

| | 24 Threads |
|----------------------|-------------------|
| Initial Write | 234.06 |
| Re-write | 311.75 |
| Random Write | 326.89 |
| Initial Read | 621.08 |
| Re-read | 662.92 |
| Random Read | 242.75 |

$$621 * 21 = 13 \text{ GByte/s}$$

Gluster v3.3



Gluster v3.4

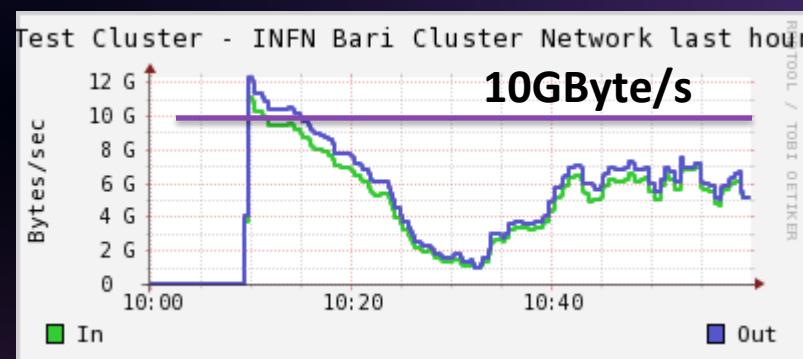
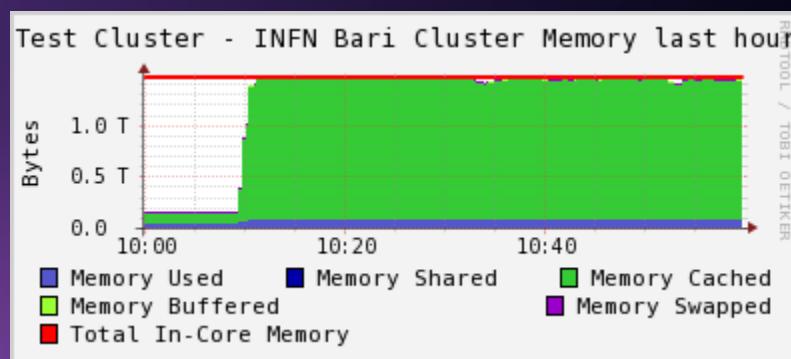
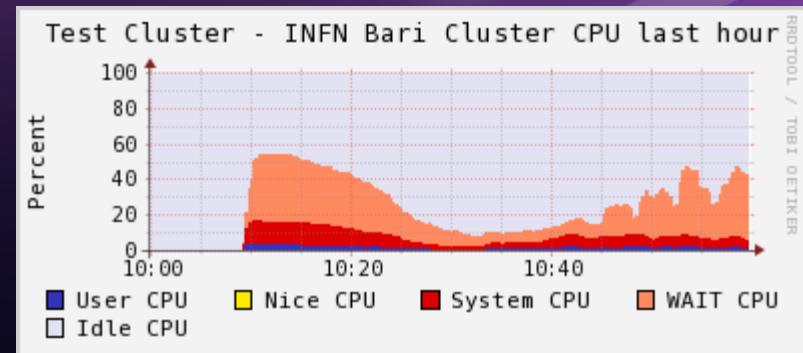
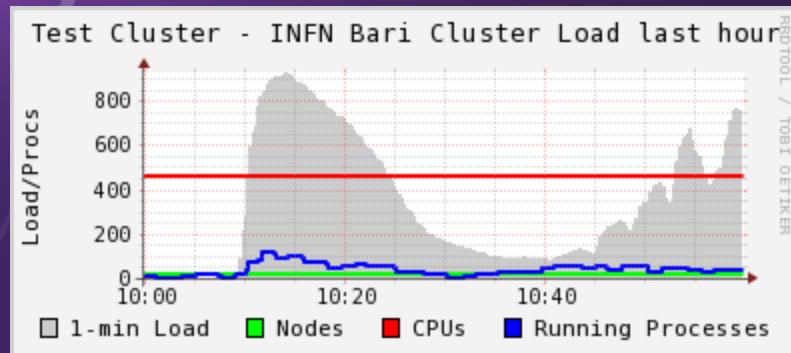
- 20 nodes, 6 brick per node
- On all the nodes (SLC6)
 - # iozone -r 128k -i 0 -i 1 -i 2 -t 24 -s 10G

MB/s

| | 24 Threads |
|---------------|------------|
| Initial Write | 306.34 |
| Re-write | 406.90 |
| Random Write | 406.33 |
| Read | 688.06 |
| Re-read | 711.46 |
| Random Read | 284.00 |

$$688 * 20 = 13 \text{ GByte/s}$$

Gluster v3.4



Using dd for comparing them all

24 dd in parallel - 10GB file - bs 4M

| MB/s | HDFS | CEPH CF | GLUSTER |
|-------|--------|---------|---------|
| read | 220.05 | 126.91 | 427.3 |
| write | 275.27 | 64.71 | 268.57 |

Average per single host (the cluster is made by 20 hosts)

Conclusions ...

- We have tested, from a point of view of the performance and functionalities, three of the main known and diffused storage solution ...
- ... trying to focus on the possibility not to use an hardware raid solution
- taking into account the new cloud storage solution that are becoming more and more interesting

Conclusions ...

- Hadoop
 - looks very stable, mature and scalable solution
 - Not fully posix compliance and not the fastest
- GlusterFS:
 - Very fast, posix compliant, and easy to manage
 - Maybe not as scalable as the others, still have few reliability problems
- CEPH:
 - Looks very scalable, complete and technological advanced
 - Still not very mature and stable, performance issues

... and future works

- We will continue this activity of testing storage solution in order to follow the quite fast evolution in this field
- In particular CEPH looks quite promising if/when stability and performance issues will be solved.
- The increasing interest in cloud storage solution are forcing the developers to put effort in providing both block and object storage solutions together with the standard posix

People Involved

- Domenico DIACONO (INFN-Bari)
- Giacinto DONVITO (INFN-Bari)
- Giovanni MARZULLI (GARR/INFN)