

Introduction to Xeon Phi

ACES

Austin, TX

Dec. 04 2013

Kent Milfeld, Luke Wilson, John
McCalpin, Lars Koesterke

TACC



THE UNIVERSITY OF TEXAS AT AUSTIN
TEXAS ADVANCED COMPUTING CENTER

What is it?


- Co-processor
 - PCI Express card
 - Stripped down Linux operating system
- Dense, simplified processor
 - Many power-hungry operations removed
 - Wider vector unit
 - Wider hardware thread count
- Lots of names
 - Many Integrated Core architecture, aka MIC
 - Knights Corner (code name)
 - Intel Xeon Phi Co-processor SE10P (product name)

What is it?

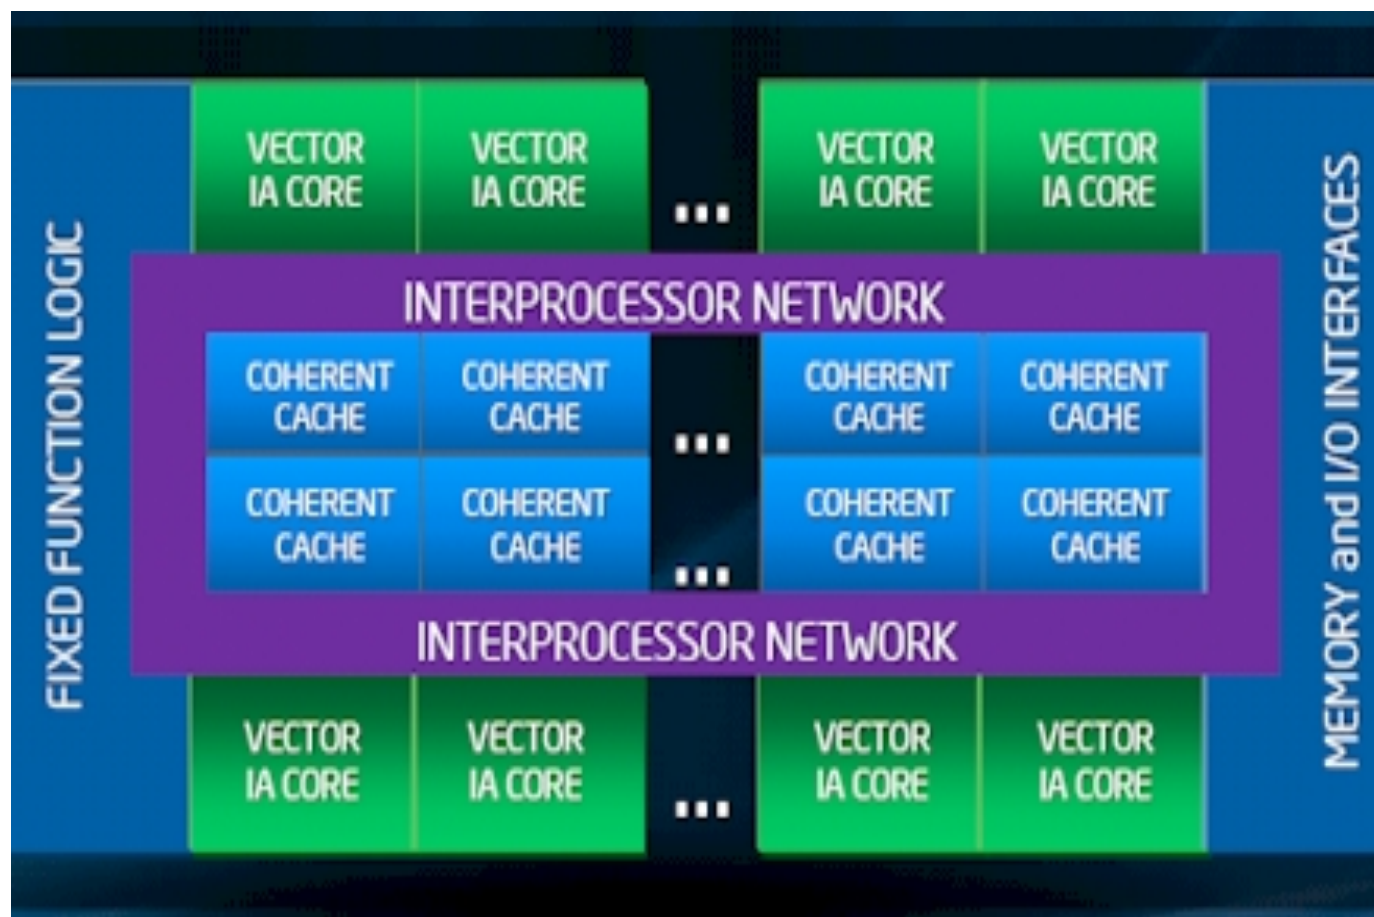
- Leverage x86 architecture (CPU with many cores)
 - x86 cores that are simpler, but allow for more compute throughput
- Leverage existing x86 programming models
- Dedicate much of the silicon to floating point ops
- Cache coherent
- Increase floating-point throughput
- Strip expensive features
 - out-of-order execution
 - branch prediction
- Widen SIMD registers for more throughput
- Fast (GDDR5) memory on card

Intel Xeon Phi Chip

- 22 nm process
- Based on what Intel learned from
 - Larrabee
 - SCC
 - TeraFlops Research Chip

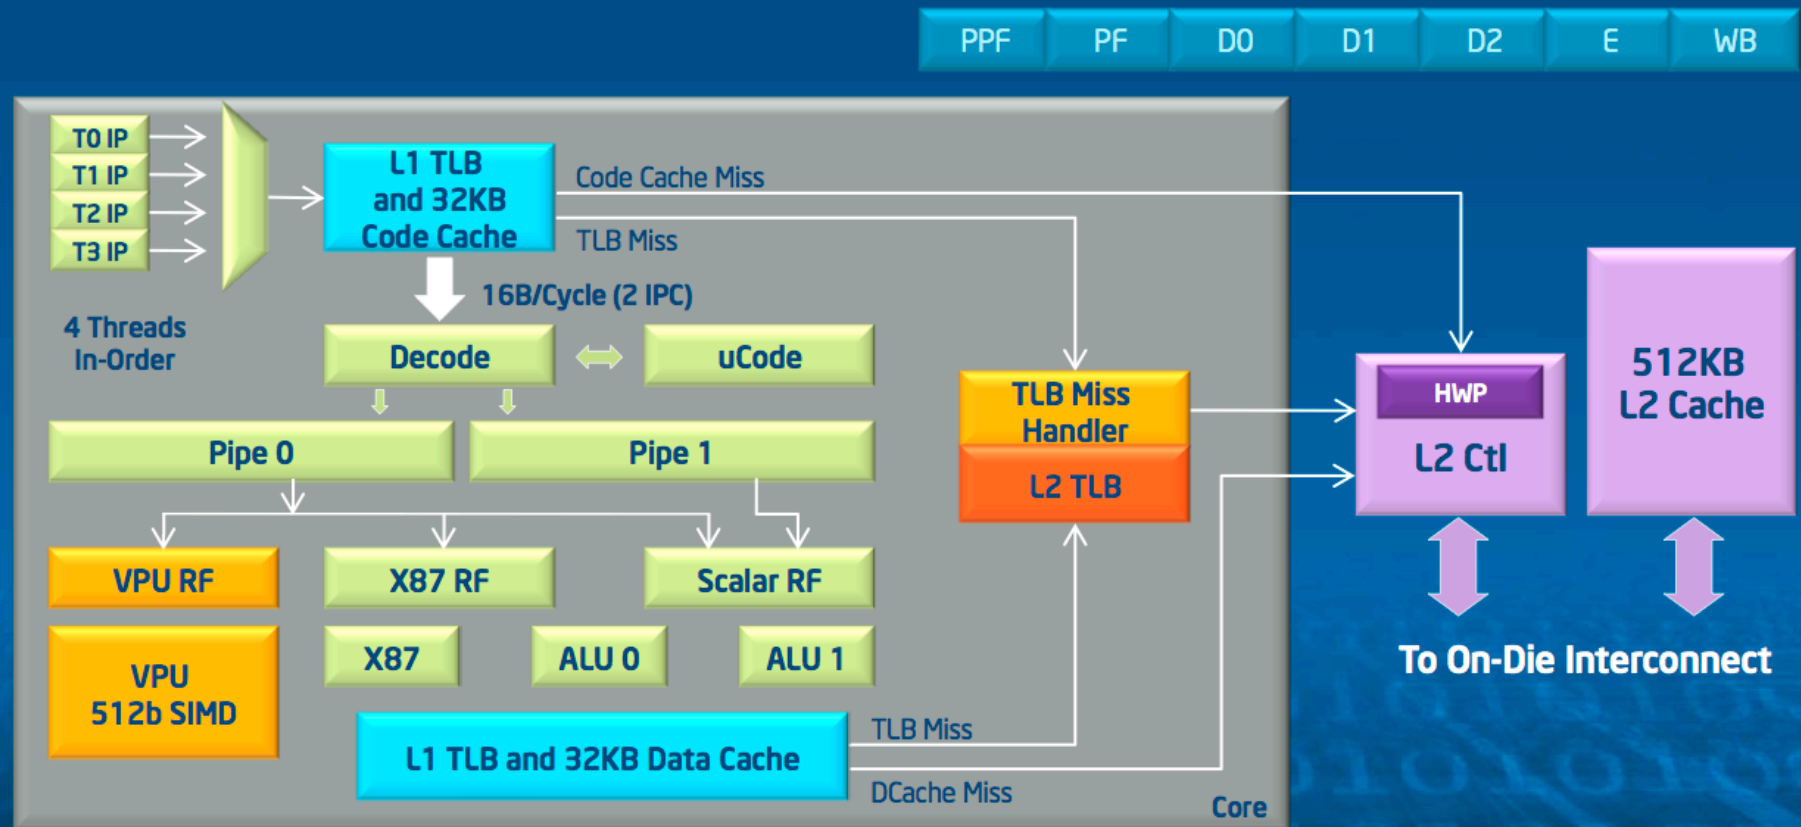


MIC Architecture



- Many cores on the die
- L1 and L2 cache
- Bidirectional ring network for L2
- Memory and PCIe connection

Knights Corner Core

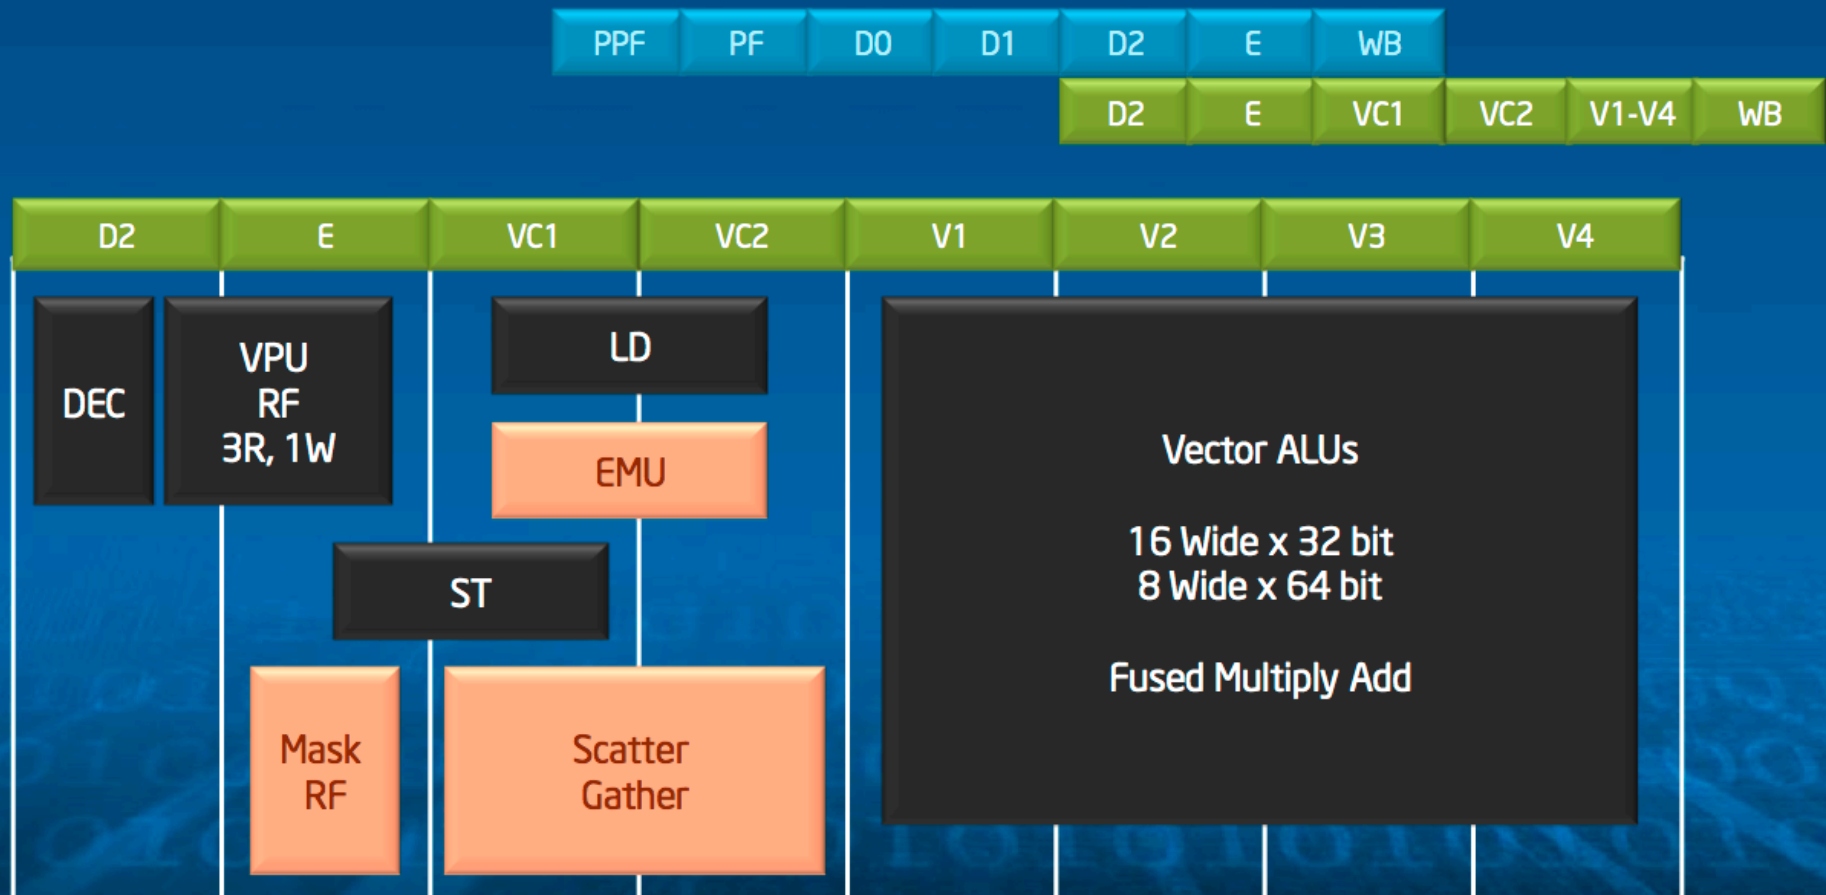


X86 specific logic < 2% of core + L2 area

George Chrysos, Intel, Hot Chips 24 (2012):

<http://www.slideshare.net/IntelXeon/under-the-armor-of-knights-corner-intel-mic-architecture-at-hotchips-2012>

Vector Processing Unit



Parallel Computing Group

Copyright © 2012 Intel Corporation. All rights reserved.

George Chrysos, Intel, Hot Chips 24 (2012):

<http://www.slideshare.net/IntelXeon/under-the-armor-of-knights-corner-intel-mic-architecture-at-hotchips-2012>

Speeds and Feeds

- Processor
 - ~1.1 GHz
 - 61 cores
 - 512-bit wide vector unit
 - 1.074 TF peak DP
- Data Cache
 - L1 32KB/core
 - L2 512KB/core, 30.5 MB/chip
- Memory
 - 8GB GDDR5 DRAM
 - 5.5 GT/s, 512-bit*
- PCIe
 - 5.0 GT/s, 16-bit

Advantages

- Intel's MIC is based on x86 technology
 - x86 cores w/ caches and cache coherency
 - SIMD instruction set
- Programming for MIC is similar to programming for CPUs
 - Familiar languages: C/C++ and Fortran
 - Familiar parallel programming models: OpenMP & MPI
 - MPI on host and on the coprocessor
 - Any code can run on MIC, not just kernels
- Optimizing for MIC is similar to optimizing for CPUs
 - **“Optimize once, run anywhere”**
 - Our early MIC porting efforts for codes “in the field” are frequently doubling performance on Sandy Bridge.

Stampede Programming Models

- Traditional Cluster
 - Pure MPI and MPI+X
 - X: OpenMP, TBB, Cilk+, OpenCL, ...
- Native Phi
 - Use one Phi and run OpenMP or MPI programs directly
- MPI tasks on Host and Phi
 - Treat the Phi (mostly) like another host
 - Pure MPI and MPI+X
- MPI on Host, Offload to Xeon Phi
 - Targeted offload through OpenMP extensions
 - Automatically offload some library routines with MKL

Traditional Cluster

- Stampede is 2+ PF of FDR-connected Xeon E5
 - High bandwidth: 56 Gb/s (sustaining >52 Gb/s)
 - Low-latency
 - $\sim 1 \mu\text{s}$ on leaf switch
 - $\sim 2.5 \mu\text{s}$ across the system
- Highly scalable for existing MPI codes
- IB multicast and collective offloads for improved collective performance

Native Execution

- Build for Phi with `-mmic`
- Execute on host
- ... or ssh to mic0 and run on the Phi
- Can safely use all 61 cores
 - Offload programs should stay away from the 61st core since the offload daemon runs here

Symmetric MPI

- Host and Phi can operate symmetrically as MPI targets
 - High code reuse
 - MPI and hybrid MPI+X
- Careful to balance workload between big cores and little cores
- Careful to create locality between local host, local Phi, remote hosts, and remote Phis
- Take advantage of topology-aware MPI interface under development in MVAPICH2
 - NSF STCI project with OSU, TACC, and SDSC

Symmetric MPI

- Typical 1-2 GB per task on the host
- Targeting 1-10 MPI tasks per Phi on Stampede
 - With 6+ threads per MPI task

MPI with Offload to Phi

- Existing codes using accelerators have already identified regions where offload works well
- Porting these to OpenMP offload should be straightforward
- Automatic offload where MKL kernel routines can be used
 - xGEMM, etc.

What we at TACC like about Phi

- Intel's MIC is based on x86 technology
 - x86 cores w/ caches and cache coherency
 - SIMD instruction set
- Programming for Phi is similar to programming for CPUs
 - Familiar languages: C/C++ and Fortran
 - Familiar parallel programming models: OpenMP & MPI
 - MPI on host and on the coprocessor
 - Any code can run on MIC, not just kernels
- Optimizing for Phi is similar to optimizing for CPUs
 - **“Optimize once, run anywhere”**
 - Our early Phi porting efforts for codes “in the field” have doubled performance on Sandy Bridge.

Will My Code Run on Xeon Phi?

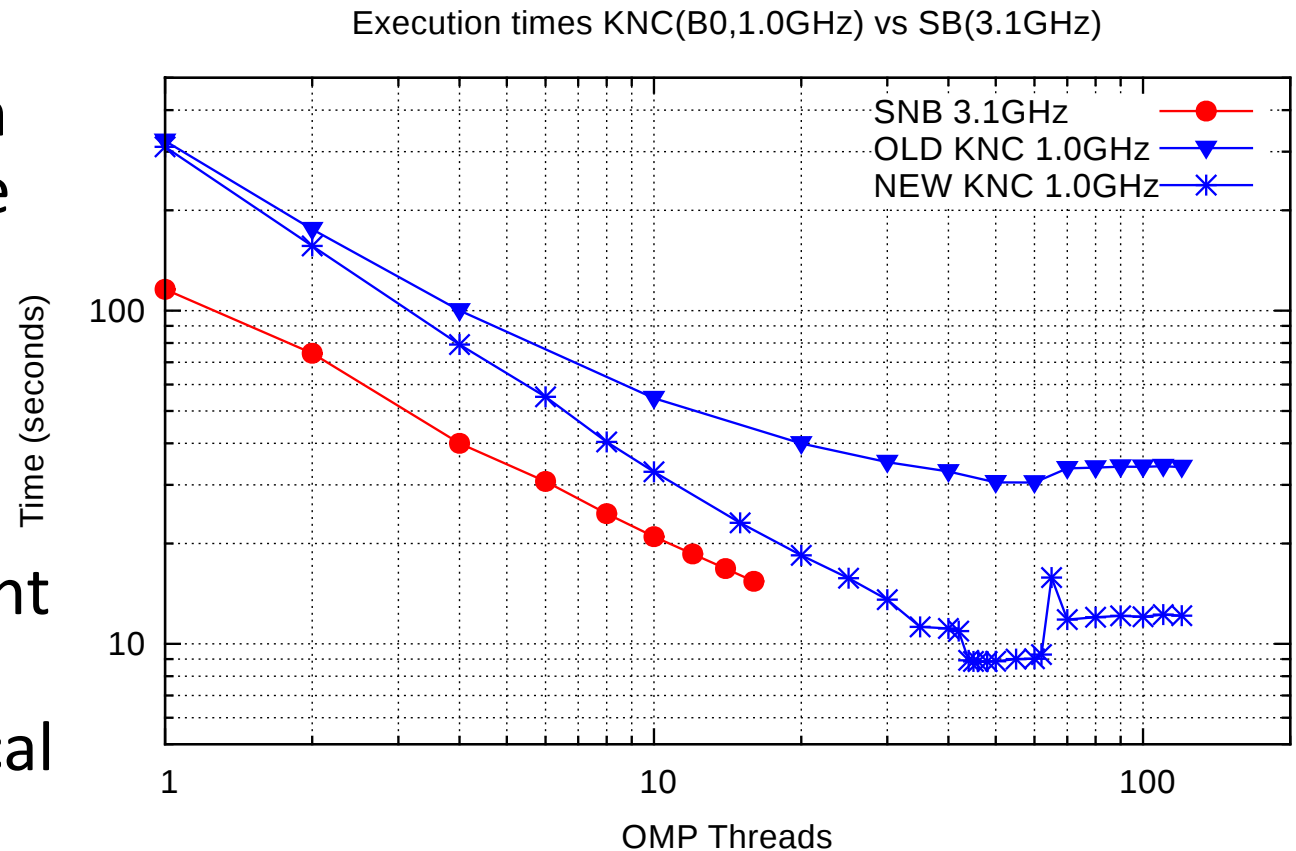
- Yes
- ... but that's the wrong question
 - Will your code run *best* on Phi?, or
 - Will you get great Phi performance without additional work?

Early Phi Programming Experiences at TACC

- Codes port easily
 - Minutes to days depending mostly on library dependencies
- Performance can require real work
 - While the software environment continues to evolve
 - Getting codes to run *at all* is almost too easy; really need to put in the effort to get what you expect
- Scalability is pretty good
 - Multiple threads per core is really important
 - Getting your code to vectorize is really important


LBM Example

- Lattice Boltzmann Method CFD code
 - Carlos Rosales, TACC
 - MPI code with OpenMP
- Finding all the right routines to parallelize is critical



PETSc/MUMPS with AO

- Hydrostatic ice sheet modeling
- MUMPS solver(called through PETSC)
- BLAS calls automatically offloaded behind the scenes



Lab I

- What you will learn
 - The lab introduces you to Stampede and to the Xeon Phi processors built into Stampede
- What you will do:
 - Compile for Xeon (host) and Xeon Phi (MIC)
 - Submit a job
 - Inspect the queue
 - Submit an interactive job
 - Execute on the host and on the Phi